

Distributed Databases

Centralised vs. Distributed Databases

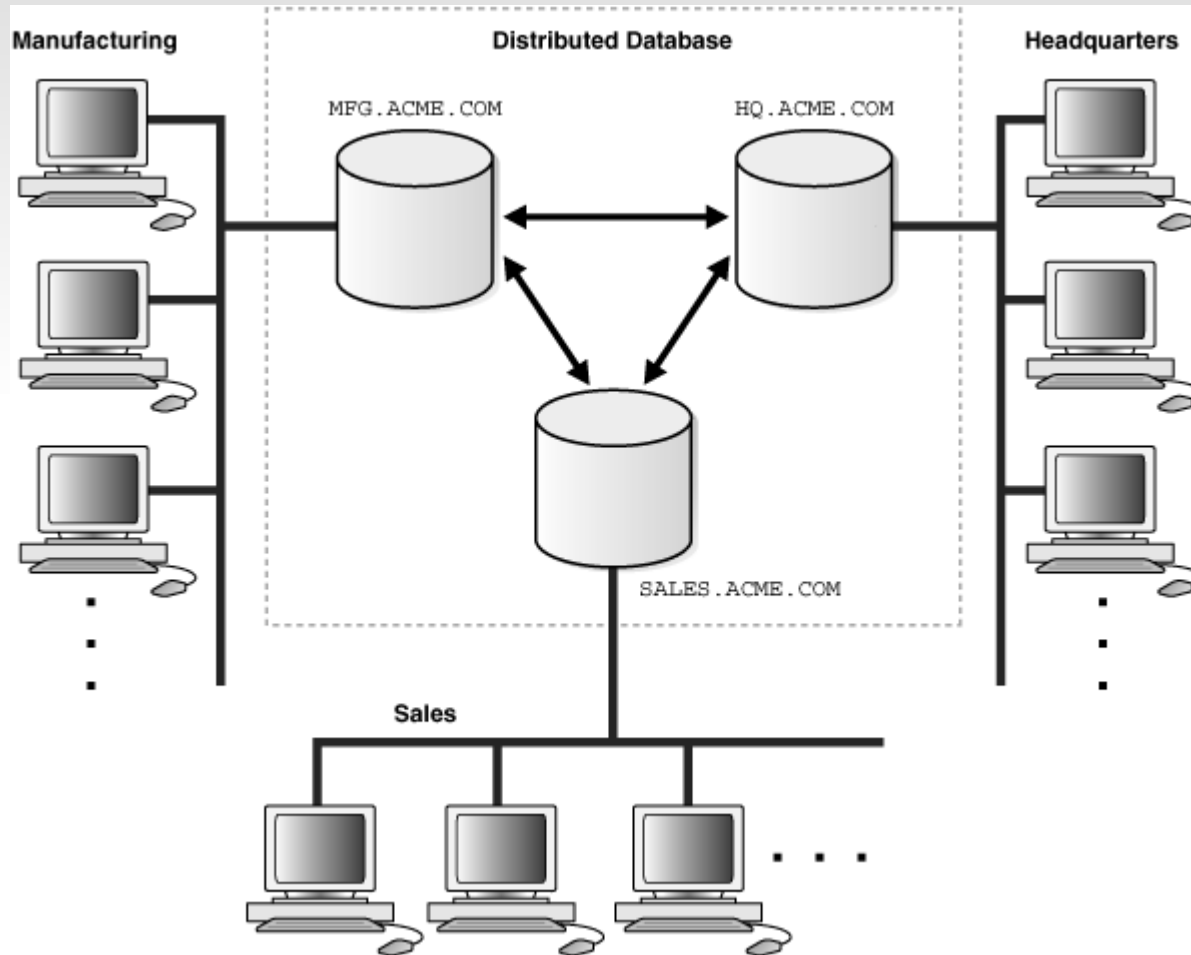
What is a Distributed Database?

- The DB is stored on several computers, interconnected through a network.
- Each site can process **local transactions** involving data only on that site.
 - Each site has a local transaction manager
 - Each site has a transaction coordinator
- Some sites may get involved in **global transactions**, which access data from several sites.
 - Local transaction managers communicate to achieve this.

Examples

- University of the West Indies
 - Bursary - Financial information
 - Personnel - Staff information
 - Registry - Student information
- Multinational
 - HQ in Kingston
 - Manufacturing in Trinidad
 - Warehouse in Miami
 - European HQ in London
 - Each site keeps data on local employees, as well as data relevant to the operation of the site.

Distributed Database Architecture Example



Why Distributed Databases I

- More natural for representing many real world organizations.
- Local autonomy
 - Local organization is fully responsible for accuracy and safety of its own portion of DB.
- Improved performance
 - Speed up of query processing
 - Possibility of parallel processing
 - Local processing
 - If data must be accessed often at a site, store it locally.

Why Distributed Databases II

- Improved availability/reliability
 - DB can continue to function even if some subsystems are down if we replicate data or hardware.
- Security
 - Avoid destruction of DB by replicating vital data.
- Incremental growth
 - Increasing size of DB
 - Increasing operations
 - Example: Include R&D facility.

What do we want from a distributed database?

- **Allow users to use the system as if it is not distributed.**
- Leads to following objectives:
 - No reliance on central *master* site.
 - Would be a bottleneck.
 - If down, whole system is down.
 - Continuous Operation
 - Enforces reliability and availability
 - Distributed Query Processing
 - A query may require accessing data at a no. of sites.
 - Same data may be available at a no. of sites.

What do we want from a distributed database?

- Distributed Transaction Management
 - Same copy of a data item may be a a number of sites.
- Transparency
- Local Autonomy

Transparency

- *Hiding* the distributed design from the user.
 - Network or distribution transparency
 - Replication transparency
 - Fragmentation transparency
 - Hardware and OS transparency
 - Language and DBMS transparency
 - Applies mostly to multidatabase systems.

Autonomy

- All operations at a site are controlled by that site.
- Types of autonomy
 - Design
 - Individual DBs can use data models and transaction management techniques that they prefer.
 - Communication
 - Individual DBs can decide which information they want to make accessible to other sites
 - Execution
 - Individual DBs can decide how to execute transactions submitted to them.

Issues and Problems I

- Distributed database design
 - How should DB and applications be placed across sites.
 - How should data dictionary be placed across sites.
 - Replication and partitioning
- Distributed query processing
 - How to break down a query into series of data manipulation operations.

Issues and Problems II

- Distributed concurrency control
 - Gets worse if data are replicated
- Distributed deadlock control
- Reliability
 - If site becomes inaccessible, how do you ensure that DBs at other sites remain consistent and up-to-date?
- Heterogeneous databases
 - Sharing pre-existing DBs with different conceptual schemas.
 - Also called multidatabase systems

Replication and Fragmentation

- Replication:
 - Should we maintain several identical copies of (part of) the DB, with each replica stored at a different site?
- Fragmentation
 - Should we partition a table into separate parts, each stored at a different site?
 - If we do, how should we partition the table?

Replication

- Advantages
 - Increased availability
 - If one site goes down, the data may be available from elsewhere.
 - Increased parallelism
 - More likely that the data needed is at the site that poses the query
 - Minimizes movement of data between sites.
 - We may send parts of the same query to different sites.

Replication

- Disadvantages
 - Increased storage space
 - Increased overhead on update.
 - Must ensure that all copies are consistent.

Fragmentation

- Why fragment? Why not simply store different complete tables at different sites?
- Two reasons:
 - Applications usually access only a subset of a relation.
 - Minimizes movement of data between sites.
 - Fragmentation gives possibility of increased parallelism.
- In fragmentation, make sure that we do not lose information.

Types of Fragmentation

- Horizontal fragmentation
 - Assign different tuples to each fragment (e.g., through a selection in the sense of relational algebra).
- Vertical fragmentation
 - Assign different attributes to each fragment.
 - Must ensure re-constructability.
- Mixed Fragmentation
 - Mixture of the two.

Problems in Fragmentation

- Increased response time if an application needs to access more than one fragment.
- Especially in vertical fragmentation, ensuring data integrity may become more difficult.
- Allocation: where to place the various fragments, and whether to replicate it.

Allocation

- Allocation concerns where to store each fragment and whether to replicate it.
- Possibilities:
 - Partitioned DB
 - Fragmentation, no replication
 - Partially replicated DB
 - Fragmentation, with each fragment stored at more than one site.
 - Fully replicated DB
 - DB is replicated in full at each site.

Evaluation of Partitioned DB

- Query processing is moderately difficult, and can be time consuming.
- Updating of DB is easy.
- Directory management is moderately difficult
- Concurrency control is easy, since can be done at one site.
- Reliability is very low.
- Requires least amount of space.

Evaluation of Partially Replicated DBs

- Query processing is moderately difficult, and can be time consuming.
- Updating of DB is moderately difficult.
- Directory management is moderately difficult.
- Concurrency control is hard.
- Reliability is high.
- Moderate amount of space.

Evaluation of Fully Replicated DBs

- Query processing is easy, and can be done quickly.
- Updating of DB is easy but time-consuming.
- Directory management is easy but time consuming.
- Concurrency control is moderately hard.
- Reliability is very high.
- Requires a lot of space.

Query Optimization

- Each DBMS has a query processor.
- The query processor takes a high level query (e.g., in SQL) and translates it into a set of relational algebraic expressions.
- Since this can be done in a number of different ways, query processor must choose the best one. This is called query optimization.

Example of query optimization

- Consider tables:
 - $\text{Empl}(\underline{\text{Eno}}, \text{Ename}, \text{Title})$
 - $\text{Job}(\underline{\text{Eno}}, \underline{\text{JobNo}}, \text{Resp}, \text{Dur})$

- Query:

```
SELECT Ename
FROM Empl, Job
WHERE Empl.Eno = Job.Eno
      AND Resp = 'Manager';
```

- Two strategies

$$\pi_{\text{ename}}(\sigma_{\text{Resp} = \text{'Manager'}}(E \bowtie J))$$

$$\pi_{\text{ename}}(E \bowtie \sigma_{\text{Resp} = \text{'Manager'}}(J))$$

Distributed Query Processing/Optimization

- Query processing/optimization is more difficult in distributed DBMS
 - Require both global optimization and local optimization
 - A query may require data from more than one site, and communication has a cost.
 - It may be possible to perform some sub-queries in parallel.
 - Cost no longer dependent on only number of tuples accessed.

Example

Site 1: $E_1 = \sigma_{\text{ENO} \leq \text{'E3'}}(E)$

Site 2: $E_2 = \sigma_{\text{ENO} > \text{'E3'}}(E)$

Site 3: $J_1 = \sigma_{\text{ENO} \leq \text{'E3'}}(J)$

Site 4: $J_2 = \sigma_{\text{ENO} > \text{'E3'}}(J)$

Results are expected at site 5.


Different Possibilities

- Strategy 1:
 - Send everything to Site 5 and perform the original query there.
- Strategy 2:
 - Do selections at sites 3 and 4
 - Send results to sites 1 and 2
 - Perform join and projections
 - Send result to site 5.
- Strategy 2 might seem better but if communication to site 5 is cheap/fast, then 1 may be better.

Distributed Transaction Management

- Ensuring ACID properties of global transactions is more complex.
- Atomicity:
 - A single transaction may involve data at many sites.
 - All sites must commit or rollback in unison.
 - Two Phase Commit can be used to enforce this.

Two-Phase Commit



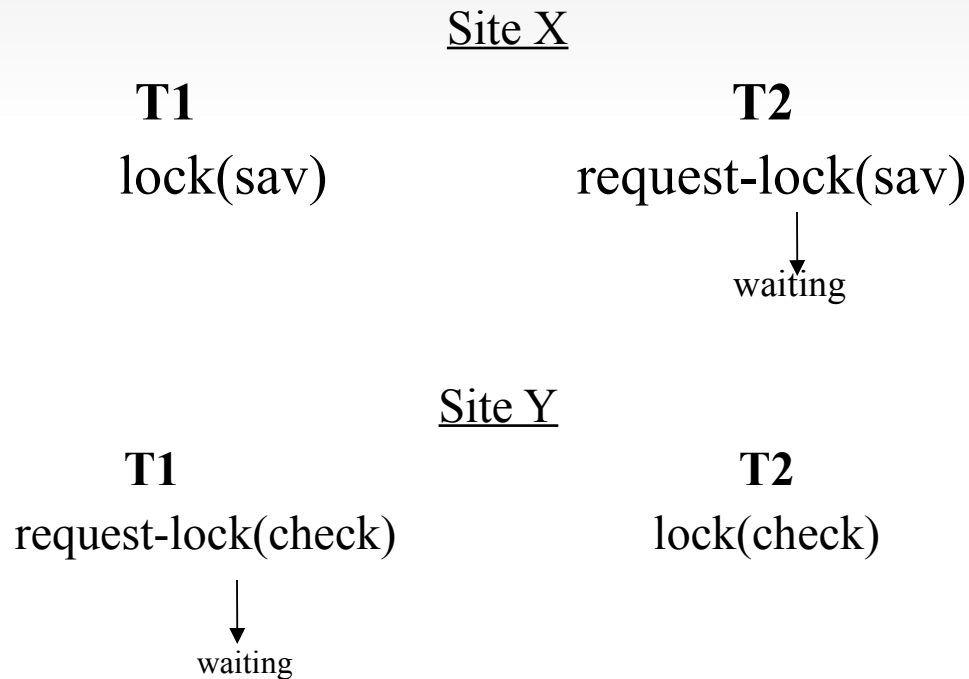
Concurrency Control

- More complex if there are multiple copies of the same data item.
- All copies of the data item must be updated (or none).
- To update an item, a lock on all copies must be granted.
- To request a lock on a copy of a data item a message must be sent to each of the sites with that copy.
 - These messages come with an overhead

- Possible to have Global Deadlocks.
 - High communication overhead to deal with this.

Global Deadlock

- E.g. Assume that data item *sav* is stored at site X and data item *check* at site Y. Both T1 and T2 need both.



Data Dictionaries

- A data dictionary in a non-distributed system contains so-called *meta-information*, information about the data.
- Examples: structure of tables, data types of attributes.
- In distributed DBMS's, data dictionary must also say where the fragments can be found.

Design of Global Data Dictionaries

- Should we have a single global directory, or a set of local ones?
- Should be maintain the data dictionary globally or in a distributed fashion?
- Should we replicate the data dictionary?

Multidatabase Systems

- Pre-existing databases form the components of multidatabase systems.
- Databases are heterogeneous.
 - Varying degrees of heterogeneity.
 - E.g. may be using different data models, dbms's, query languages, naming conventions.
- Provide logical data integration without requiring physical data integration.
 - Must address the heterogeneity issues.
 - Even more difficult to design than distributed systems.
 - Wrapper-Mediator approach is commonly used.